

## [Faster Access to CILogon's Persistent Store](#)

[Motivation](#)

[Proposed solution](#)

[Security and access](#)

[The API itself](#)

[Tomcat Advisory](#)

[Making a request](#)

[Format of a response from the server](#)

[Status Codes and Error Conditions](#)

[Date formatting](#)

[The API Calls](#)

[GetUserID\(remoteUser, idp\)](#)

[GetUser](#)

[Case #1 - getUser\(uid\)](#)

[Case #2 - getUser\(remoteUser, idp,...\)](#)

[RemoveUser\(uid\)](#)

[Set all IDPs](#)

[Get All IDPS](#)

[Get Portal Parameters](#)

[GetLastArchivedUser](#)

[setTwoFactorInfo\(userid, twoFactorInfo\)](#)

[getTwoFactorInfo\(userid\)](#)

## **Faster Access to CILogon's Persistent Store**

### **Motivation**

Current access is through the PHP layer which in turn makes a call to a Perl database application. This is causing some very noticeable performance issues. The main reasons are

- The PHP library creates a completely new Perl interpreter for each call. This has a fixed overhead of .2 sec
- The Perl application requires roughly .5 seconds to create a connection to the database.

No connection pooling is possible because of issue #1, so regardless of the database's speed, there is a fixed .7 second delay for *each* call to the database. The net effect is an unacceptably slow web application.

### **Proposed solution**

The Java delegation application manages all of its state with connection pooling and caches. Although there is a small initial cost for initialization whenever Tomcat is restarted (< 3 sec.) , access afterwards runs consistently in 20ms - 30 ms range. This is an improvement over the current method by an easy factor of 10. It makes sense to expose this to the rest of the CILogon Service itself. Under Tomcat, each web application with its various servlets share a common instance. Rather than set up a separate dedicated web application to serve this, it will be a simple servlet running under the delegation application. The main reason for this is that otherwise there will be both threading and latency issues involved in updating and accessing information which would be very difficult to resolve.

## Security and access

For performance reasons, access will be via http (no SSL) to Tomcat directly on port 8080 (avoiding the translation from Apache through the AJP connector). Since this in effect would make *everything* in the database world readable, only localhost access will be permitted. The path to the database servlet will be **http://localhost:8080/delegation/dbService**. This is the standard and the strongly suggested way to do this.

Alternately, If you are running it in another environment, you should enable SSL and may set DB Service users in the configuration file.

```
<dbService enabled="true">
  <user>
    <user name="..." hash="..." />
    <!-- more users -->
  </users>
</dbService>
```

In this case, the user has a name and an issued password. The SHA-1 hash of the password is recorded here. When calling the DB Service, the user name and password must be supplied with each call. the default is to have this disabled (so if this section is missing from your configuration file it is disabled) since the suggested way to do this is by not having a publicly accessible Tomcat.

## The API itself

All access will be via HTTP Get. Technically we should only allow HTTP Post on requests which may result in changes to the server (e.g. a creating a user) and HTTP Get for information requests, however, this would make it awkward to use. The aim is to make it simple enough that clients can hand-code the request and parse the result without having to resort to a library or other mechanism to do so.

## Tomcat Advisory

(Note applies to Tomcat 5 only.) Allowing access via HTTP GET does have a potential limit.

Tomcat will refuse to accept any request over a default size of 8192 bytes, so in particular setting IDPs might fail at some point. The fix is to set the `maxHttpHeaderSize` parameter in the connector e.g., as follows

```
<Connector port="8080"
  protocol="HTTP/1.1"
  maxHttpHeaderSize="65535"
  connectionTimeout="20000"
  redirectPort="8443" />
```

Failure to do so will lead, at some point to (misleading) HTTP 404 errors.

## Making a request

Every request will go to the same address:

**`http://localhost:8080/delegation/dbService.`**

Requests are standard HTTP GET with key=value pairs. One of these is required, "action=XXXX", where XXXX determines what is to be done by the server. The remaining key/value pairs are parameters for the call. All arguments in a post or get will be URL encoded but key/value pairs may be in any order. Required arguments must be present or an error will be returned. Generally (except for a list of IDPs) repeated arguments will cause a duplicate argument message to be generated.

## Format of a response from the server

The general serialized form for an object (which is always the body of the response and is urlencoded) is:

```
status=XXX
key1=value1
key2=value2
...
```

where keys may be repeated for multi-valued objects. Each key/value pair is followed by a linefeed, except the final one. If a value is missing, it will still be present in the form "key=" (no value). Each section below detailing an API call will list what are and are not acceptable values. The status line is always present. The next section details the possible values.

## Status Codes and Error Conditions

It may occur that there are errors during the processing of a request, e.g. if the parameters are incorrect. Rather than confuse server errors with data errors, we follow the following policy: Only actual errors with the servlet itself will result in an HTTP status code different than 200. E.g. a 404 Not Found error can only have come from the web server itself and means there is a problem with the request rather than meaning, say, that a user was not found. The first line of each response (and the only required line) is the status. This will be recorded in the body of the response according to the following table. All operations except saving the list of IDPs can return a duplicate parameter exception. It is not the task of this servlet to disambiguate or merge conflicting requests. Duplicate arguments are in general not allowed for any argument, including optional ones. The exception to this rule is setting a list of IDPs. All operations can return a

missing parameter error for a required parameter.

| Value                   | Decimal | Hex     | Comment                                                                                                      |
|-------------------------|---------|---------|--------------------------------------------------------------------------------------------------------------|
| OK                      | 0       | 0x0     | Normal return                                                                                                |
| ActionNotFound          | 1       | 0x1     | No such action is supported by this service. Normally this indicates that the action value was mis-typed.    |
| NewUser                 | 2       | 0x2     | A new user was created                                                                                       |
| UserUpdated             | 4       | 0x4     | An existing user was updated                                                                                 |
| UserNotFound            | 6       | 0x6     | The requested user was not found. Returned by hasUser, getLastArchivedUser, removeUser.                      |
| UserExists              | 8       | 0x8     | The requested user exists. Returned by hasUser.                                                              |
| UserExistsError         | 1048481 | 0xFFFA1 | An attempt to get a nonexistent user failed.                                                                 |
| UserNotFoundError       | 1048483 | 0xFFFA3 | An attempt to update or otherwise access a nonexistent user failed.                                          |
| TransactionNotFound     | 1048485 | 0xFFFA5 | No such transaction for the given temporary credential was found                                             |
| IDPSaveFailed           | 1048487 | 0xFFFA7 | Saving the list of idps failed.                                                                              |
| DuplicateParameterFound | 1048561 | 0xFFFF1 | A duplicate argument was supplied.                                                                           |
| InternalError           | 1048563 | 0xFFFF3 | Some error internal to the server occurred during processing. Consult the server logs.                       |
| SaveIDPFailed           | 1048565 | 0xFFFF5 | There was a problem saving the list of IDPs.                                                                 |
| MalformedInputError     | 1048567 | 0xFFFF7 | An input was of the incorrect format. E.g. an illegal uri or a string that cannot be parsed into an integer. |
| MissingParameterError   | 1048569 | 0xFFFF9 | A required argument was not found.                                                                           |
| NoRemoteUser            | 1048571 | 0xFFFFB | Calls that require the remote user will fail with this message if it is not supplied.                        |
| NoIdentityProvider      | 1048573 | 0xFFFFD | Calls that require the identity provider will                                                                |

|                           |       |       |                                                                                            |
|---------------------------|-------|-------|--------------------------------------------------------------------------------------------|
|                           |       |       | fail with this message if it is not supplied.                                              |
| Transaction not found     | 65537 | 10001 | No transaction with the given identifier could be found                                    |
| Expired token             | 65539 | 10003 | The token for this request has expired.                                                    |
| Create transaction failed | 65541 | 10005 | General exception when a transaction cannot be created                                     |
| Unknown callback          | 65543 | 10007 | The supplied callback id not in the list of registered callbacks                           |
| Client id missing         | 65545 | 10009 | No client identifier has been supplied with this request                                   |
| No registered callbacks   | 65547 | 1000A | The client has no callbacks registered. (Normally this implies an incomplete registration) |
| Unknown client            | 65549 | 1000C | The identifier does not match any client                                                   |
| Unapproved client         | 65551 | 1000E | This client has been registered but has not yet been approved.                             |

## Date formatting

As with the rest of the system, all dates are ISO 8601 compliant of the form

yyyy-mm-ddThh:mm:ss.xxxxZ

Where the infix T prefixes the time and the final Z indicates this is GMT. E.g.,

2010-11-17T17:09:19.692Z

Dates are needed for instance on user and archived user objects.

# The API Calls

## GetUserID(remoteUser | EPPN | EPTID | OpenID|oidc, idp)

**What's it do:** Given the at least one of the identifiers and the idp, this will return the internal unique identifier for this user.

**Examples:**

[http://localhost:8080/delegation/dbService?action=getUserID&remote\\_user=bob%40foo.edu&idp=urn%3Aincommon%3Auiuc.edu](http://localhost:8080/delegation/dbService?action=getUserID&remote_user=bob%40foo.edu&idp=urn%3Aincommon%3Auiuc.edu)

<http://localhost:8080/delegation/dbService?action=getUserID&eppn=bob%40foo.edu&eptid=https%3A%2F%2Fidp.uni.edu.au%2Fidp%2Fshibboleth!https%3A%2F>

[http://localhost:8080/delegation/dbService?  
action=getUserID&open\\_id=bob2468%40myopen.com&idp=urn%3Aaaf.edu](http://localhost:8080/delegation/dbService?action=getUserID&open_id=bob2468%40myopen.com&idp=urn%3Aaaf.edu)

[http://localhost:8080/delegation/dbService?  
action=getUserID&open\\_id=bob2468%40myopen.com&idp=urn%3Aaaf.edu](http://localhost:8080/delegation/dbService?action=getUserID&open_id=bob2468%40myopen.com&idp=urn%3Aaaf.edu)

Request key/value pairs.

| Key         | Value                      | Comment                                                                        |
|-------------|----------------------------|--------------------------------------------------------------------------------|
| action      | getUserID                  | Required. Get the unique identifier for the user given the remote user and idp |
| remote_user | remote user                | Required*                                                                      |
| eppn        | EduPerson Principal Name   | Required*                                                                      |
| eptid       | EduPerson Targeted ID      | Required*                                                                      |
| open_id     | Open ID                    | Required*                                                                      |
| oidc        | Open ID Connect identifier | Required*                                                                      |
| idp         | the identity provider      | Required.                                                                      |

\* = at least one of the values must be present.

Response key/values pairs

| Key      | Value                            | Comment |
|----------|----------------------------------|---------|
| status   | OK, UserNotFound, ActionNotFound |         |
| user_uid | the unique identifier            |         |

A complete response, for example:

HTTP/1.1 200 OK

Server: Apache-Coyote/1.1

Content-Type: application/x-www-form-urlencoded

Transfer-Encoding: chunked

Date: Fri, 19 Nov 2010 22:51:55 GMT

48

status=OK

user\_uid=https%3A%2F%2Fci.ilogon.org%2FserverA%2Fusers%2F1803

# GetUser

**What's it do:** This set of calls will always return the complete set of user information as stored. Missing values will be returned as either "key=" with no value or omitted completely

Note: This has 2 different possible polymorphic calls associated. We will give each in terms of its arguments separately

## Case #1 - getUser(uid)

**Example:**

**Request:**

http://localhost:8080/delegation/dbService?action=getUser&user\_uid=http%3A%2F%2Fcilogon.org%2FserverA%2Fusers%2F119

**Body of response:**

status=OK  
remote\_user=gaynor%40illinois.edu  
idp=urn%3Aamace%3Aincommon%3Auiuc.edu  
idp\_display\_name=University+of+Illinois+at+Urbana-Champaign  
first\_name=Jeffrey  
last\_name=Gaynor  
user\_uid=http%3A%2F%2Fcilogon.org%2FserverA%2Fusers%2F119  
email=gaynor%40illinois.edu  
serial\_string=A119  
distinguished\_name=%2FDC%3Dorg%2FDC%3Dcilogon%2FC%3DUS%2F0%3DUniversity+of+Illinois+at+Urbana-Champaign%2FCN%3DJeffrey+Gaynor+A119+email%3Dgaynor%40illinois.edu  
create\_time=2010-05-14T22%3A23%3A23.397Z  
two\_factor=sjkhsdkkfhsirandomStringishThingish

Request key/value pairs.

| Key      | Value                     | Comment   |
|----------|---------------------------|-----------|
| action   | getUser                   | Required. |
| user_uid | the unique id of the user | Required. |

Response key/value pairs

| Key         | Value                                                                               | Comment                   |
|-------------|-------------------------------------------------------------------------------------|---------------------------|
| status      | OK, ActionNotFound<br>UserNotFoundError,<br>MissingParameter,<br>DuplicateParameter |                           |
| user_uid    |                                                                                     | Identical to the argument |
| remote_user |                                                                                     |                           |

|                    |                                                         |                                                                              |
|--------------------|---------------------------------------------------------|------------------------------------------------------------------------------|
| idp                |                                                         |                                                                              |
| idp_display_name   |                                                         |                                                                              |
| first_name         |                                                         |                                                                              |
| last_name          |                                                         |                                                                              |
| email              |                                                         |                                                                              |
| serial_string      | The serial string (e.g. A123) as computed by the server |                                                                              |
| distinguished_name | The DN as computed by the server                        |                                                                              |
| create_time        |                                                         | ISO 8601 formatted date field.                                               |
| two_factor         |                                                         | A string with two factor information if available. Omitted if none is found. |

## Case #2 - getUser(remoteUser|EPPN|EPTID|OpenID|oidc, idp,...)

**Note:** This version of getUser is actually a shortcut for one of

- creating a new user,
- fetching a user with given identifier and idp fields
- checking such an existing user, updating its fields and returning the result.

The identifier and idp are the only required arguments. Omitting the others will set them to being empty, so all 6 parameters are always required. This means in particular this should be viewed as an update to a user which happens to return the user too. If the user has changed, then it is automatically archived, the most recent of which may be recovered with the getLastArchivedUser call detailed below.

Request key/value pairs

| Key         | Value   | Comment   |
|-------------|---------|-----------|
| action      | getUser | Required. |
| remote_user |         | Required* |
| eppn        |         | Required* |
| eptid       |         | Required* |
| open_id     |         | Required* |
| oidc        |         | Required* |



|                  |  |           |
|------------------|--|-----------|
| idp              |  | Required. |
| idp_display_name |  |           |
| first_name       |  |           |
| last_name        |  |           |
| email            |  |           |

\* = at least one of these must be included.

#### Response key-value pairs

| Key              | Value                                                                          | Comment                    |
|------------------|--------------------------------------------------------------------------------|----------------------------|
| status           | OK,ActionNotFound,<br>UserNotFound,<br>MissingParameter,<br>DuplicateParameter |                            |
| user_uid         |                                                                                |                            |
| remote_user      |                                                                                | Identical to the argument* |
| eppn             |                                                                                | Identical to the argument* |
| eptid            |                                                                                | Identical to the argument* |
| open_id          |                                                                                | Identical to the argument* |
| oidc             |                                                                                | identical to the argument* |
| idp              |                                                                                | Identical to the argument  |
| idp_display_name |                                                                                | Identical to the argument  |
| first_name       |                                                                                | Identical to the argument  |
| last_name        |                                                                                | Identical to the argument  |
| email            |                                                                                | Identical to the argument  |
| serial_string    | See above                                                                      |                            |
| dn               | See above                                                                      |                            |

\* = **IF** this is supplied as an argument, it will be identical. It is possible that other fields may be returned as well. E.g. If both eppn and eptid are stored, you may request a user by their eptid only. Then you would get that plus the stored eppn.

|            |  |                                                             |
|------------|--|-------------------------------------------------------------|
| two_factor |  | Two factor information, if any is found. Omitted otherwise. |
|            |  |                                                             |

## RemoveUser(uid)

**What's it do:** Removes the user from the store, archiving it beforehand.

**Example:**

**Request:** [http://localhost:8080/delegation/dbService?action=removeUser&user\\_uid=https%3a%2f%2fcilogon.org%2fserverA%2fusers%2f2133](http://localhost:8080/delegation/dbService?action=removeUser&user_uid=https%3a%2f%2fcilogon.org%2fserverA%2fusers%2f2133)

**Body of response:**

status=0

Request key/value pairs

| Key      | Value                        | Comment   |
|----------|------------------------------|-----------|
| action   | removeUser                   | Required. |
| user_uid | The user's unique identifier | Required  |

Response key/value pairs

| Key    | Value                                                                       | Comment                                                         |
|--------|-----------------------------------------------------------------------------|-----------------------------------------------------------------|
| status | OK, ActionNotFound, MissingParameter, DuplicateParameter, UserNotFoundError | If the asserted user is not in the store, an error is returned. |

## Set all IDPs

**What's it do:** Save the entire given list. This replaces the current list of IDPs.

**Note:** This will not save an empty list of IDPs. If there is an error, this will try to rollback the save.

**Example:** (With url encoding)

[http://localhost:8080/delegation/dbService?action=setAllIdps&idp\\_uid=urn%253Aidentity%252Fprov%252F1290201592400&idp\\_uid=urn%253Aidentity%252Fprov%252F1290201595564](http://localhost:8080/delegation/dbService?action=setAllIdps&idp_uid=urn%253Aidentity%252Fprov%252F1290201592400&idp_uid=urn%253Aidentity%252Fprov%252F1290201595564)

Request key/value pairs

| Key     | Value                           | Comment                      |
|---------|---------------------------------|------------------------------|
| action  | setAllIdps                      | Required.                    |
| idp_uid | A url encoded identity provider | Required (multiples allowed) |

Response key/value pairs

| Key    | Value                                | Comment                                                       |
|--------|--------------------------------------|---------------------------------------------------------------|
| status | OK, ActionNotFound, MissingParameter | A missing parameter exception occurs if no idps are supplied. |

## Get All IDPS

**What's it do:** Get the list of all current IDPs

**Example:**

**Request:** `http://localhost:8080/dbService?action=getAllIdps`

**Body of Response:**

status=OK

idp\_uid=urn%3Apace%3Aincommon%3Auchicago.edu

idp\_uid=urn%3Apace%3Aincommon%3Aidp.protectnetwork.org

idp\_uid=urn%3Apace%3Aincommon%3Aibl.gov

Request key/value pairs

| Key    | Value      | Comment   |
|--------|------------|-----------|
| action | getAllIdps | Required. |

Response key/value pairs

| Key     | Value                           | Comment                                             |
|---------|---------------------------------|-----------------------------------------------------|
| status  | OK, ActionNotFound              |                                                     |
| idp_uid | a url encoded identity provider | There will be many of these normally, one per line. |

## Get Portal Parameters

**What's it do:** Gets the portal's parameters (callback, name, etc.) based on the temporary credential.

Request key/value pairs

| Key         | Value                    | Comment   |
|-------------|--------------------------|-----------|
| action      | getPortalParameter       | Required. |
| oauth_token | The temporary credential | Required. |

Response key/value pairs

| Key                 | Value                                                                         | Comment          |
|---------------------|-------------------------------------------------------------------------------|------------------|
| status              | OK, ActionNotFound, TransactionNotFound, MissingParameter, DuplicateParameter |                  |
| oauth_token         |                                                                               | same as argument |
| cilogon_callback    | The callback uri                                                              |                  |
| cilogon_success     | The success uri                                                               |                  |
| cilogon_failure     | The failure uri                                                               |                  |
| cilogon_portal_name | The name of the portal                                                        |                  |

## GetLastArchivedUser

**What's it do:** This will take the id of a user and return the last user archived under that id, if there is one.

**Note:** The internal information about the archive (the date at which the user was archived) is not returned. Just a user object. Also, users are archived as a matter of course when during the getUser call, if the user's information has changed. See the information under that entry.

Request key/value pairs

| Key      | Value                        | Comment   |
|----------|------------------------------|-----------|
| action   | getLastArchivedUser          | Required. |
| user_uid | The user's unique identifier | Required  |

Response key/value pairs

| Key         | Value                                                     | Comment |
|-------------|-----------------------------------------------------------|---------|
| status      | OK, DuplicateParameter, ActionNotFound, UserNotFoundError |         |
| user_uid    |                                                           |         |
| remote_user |                                                           |         |
| eppn        |                                                           |         |
| eptid       |                                                           |         |

|                  |  |  |
|------------------|--|--|
| open_id          |  |  |
| oidc             |  |  |
| idp              |  |  |
| idp_display_name |  |  |
| first_name       |  |  |
| last_name        |  |  |
| email            |  |  |
| serial_string    |  |  |
| dn               |  |  |

## setTwoFactorInfo(userid, twoFactorInfo)

**What's it do?** Set the two factor information for the given user id.

**Note:** The two factor information is simply an opaque string. No parsing or other processing of it will be done.

Request key/value pairs

| Key        | Value            | Comment                                                                          |
|------------|------------------|----------------------------------------------------------------------------------|
| action     | setTwoFactorInfo |                                                                                  |
| user_uid   | user identifier  |                                                                                  |
| two_factor | String           | Opaque string. No processing of any sort done. This is not assumed to be binary. |

Response key/value pairs

| Key    | Value              | Comment |
|--------|--------------------|---------|
| status | ok, user not found |         |

## getTwoFactorInfo(userid)

**What's it do?** This will retrieve any two factor information from the server for the given user id.

Request key/value pairs

| Key      | Value                  | Comment |
|----------|------------------------|---------|
| action   | getTwoFactorInfo       |         |
| user_uid | user unique identifier |         |

Response key/value pairs

| Key        | Value              | Comment                                                                                     |
|------------|--------------------|---------------------------------------------------------------------------------------------|
| status     | ok, user not found |                                                                                             |
| user_uid   | user id            | Same as in the request.                                                                     |
| two_factor | string             | as per setter, this is assumed to be an opaque string. Omitted if no user information found |

## setTransactionState()

**What's it do?** (OAuth2) This will retrieve the transaction corresponding to a given authorization grant.

## createTransaction

**What's it do?** (OAuth2) This creates a transaction. Note that the way CILogon works, the call goes to a PHP layer first which in turn has to do various bits of setup. This call is made by the PHP layer to set up an OAuth transaction and is the first call of any sort after a request is received.

## getClient(client\_id)

**What's it do?** (OAuth2) This will retrieve the information for a given client.